

# Lecture 13: Fast Reinforcement Learning<sup>1</sup>

Emma Brunskill

CS234 Reinforcement Learning


Winter 2020

---

<sup>1</sup>With a few slides derived from David Silver

# Refresh Your Knowledge Fast RL Part II

- The prior over arm 1 is Beta(1,2) (left) and arm 2 is a Beta(1,1) (right figure). Select all that are true.
  - 1 Sample 3 params: 0.1,0.5,0.3. These are more likely to come from the Beta(1,2) distribution than Beta(1,1).
  - 2 Sample 3 params: 0.2,0.5,0.8. These are more likely to come from the Beta(1,1) distribution than Beta(1,2).
  - 3 It is impossible that the true Bernoulli param is 0 if the prior is Beta(1,1).
  - 4 Not sure
- The prior over arm 1 is Beta(1,2) (left) and arm 2 is a Beta(1,1) (right). The true parameters are arm 1  $\theta_1 = 0.4$  & arm 2  $\theta_2 = 0.6$ . Thompson sampling = TS
  - 1 TS could sample  $\theta = 0.5$  (arm 1) and  $\theta = 0.55$  (arm 2).
  - 2 For the sampled thetas (0.5,0.55) TS is optimistic with respect to the true arm parameters for all arms.
  - 3 For the sampled thetas (0.5,0.55) TS will choose the true optimal arm for this round.
  - 4 Not sure



# Class Structure

- Last time: Fast Learning (Bayesian bandits to MDPs)
- **This time: Fast Learning III (MDPs)**
- Next time: Batch RL

# Settings, Frameworks & Approaches

- Over these 3 lectures will consider 2 settings, multiple frameworks, and approaches
- Settings: Bandits (single decisions), MDPs
- Frameworks: evaluation criteria for formally assessing the quality of a RL algorithm. So far seen empirical evaluations, asymptotic convergence, regret, probably approximately correct
- Approaches: Classes of algorithms for achieving particular evaluation criteria in a certain set. So far for exploration seen: greedy,  $\epsilon$ -greedy, optimism, Thompson sampling, for multi-armed bandits

# Table of Contents

1 MDPs

2 Bayesian MDPs

3 Generalization and Exploration

4 Summary

# Fast RL in Markov Decision Processes

- Very similar set of frameworks and approaches are relevant for fast learning in reinforcement learning
- Frameworks
  - Regret
  - Bayesian regret
  - Probably approximately correct (PAC)
- Approaches
  - Optimism under uncertainty
  - Probability matching / Thompson sampling
- Framework: Probably approximately correct

# Fast RL in Markov Decision Processes

- Montezuma's revenge
- [https://www.youtube.com/watch?v=ToSe\\_CUG0F4](https://www.youtube.com/watch?v=ToSe_CUG0F4)

# Model-Based Interval Estimation with Exploration Bonus (MBIE-EB)

(Strehl and Littman, J of Computer & Sciences 2008)

- 
- 1: Given  $\epsilon, \delta, m$
  - 2:  $\beta = \frac{1}{1-\gamma} \sqrt{0.5 \ln(2|S||A|m/\delta)}$
  - 3:  $n_{sas}(s, a, s') = 0, \forall s \in S, a \in A, s' \in S$
  - 4:  $rc(s, a) = 0, n_{sa}(s, a) = 0, \tilde{Q}(s, a) = 1/(1 - \gamma), \forall s \in S, a \in A$
  - 5:  $t = 0, s_t = s_{init}$
  - 6: **loop**
  - 7:    $a_t = \arg \max_{a \in A} \tilde{Q}(s_t, a)$
  - 8:   Observe reward  $r_t$  and state  $s_{t+1}$
  - 9:    $n_{sa}(s_t, a_t) = n_{sa}(s_t, a_t) + 1, n_{sas}(s_t, a_t, s_{t+1}) = n_{sas}(s_t, a_t, s_{t+1}) + 1$
  - 10:    $rc(s_t, a_t) = \frac{rc(s_t, a_t)(n_{sa}(s_t, a_t) - 1) + r_t}{n_{sa}(s_t, a_t)}$
  - 11:    $\hat{R}(s_t, a_t) = rc(s_t, a_t)$  and  $\hat{T}(s'|s_t, a_t) = \frac{n_{sas}(s_t, a_t, s')}{n_{sa}(s_t, a_t)}, \forall s' \in S$
  - 12:   **while** not converged **do**
  - 13:      $\tilde{Q}(s, a) = \hat{R}(s, a) + \gamma \sum_{s'} \hat{T}(s'|s, a) \max_{a'} \tilde{Q}(s', a) + \frac{\beta}{\sqrt{n_{sa}(s, a)}}, \forall s \in S, a \in A$
  - 14:   **end while**
  - 15: **end loop**



# Framework: PAC for MDPs

- For a given  $\epsilon$  and  $\delta$ , A RL algorithm  $\mathcal{A}$  is PAC if on all but  $N$  steps, the action selected by algorithm  $\mathcal{A}$  on time step  $t$ ,  $a_t$ , is  $\epsilon$ -close to the optimal action, where  $N$  is a polynomial function of  $(|S|, |A|, \gamma, \epsilon, \delta)$
- Is this true for all algorithms?

# MBIE-EB is a PAC RL Algorithm

**Theorem 2.** Suppose that  $\epsilon$  and  $\delta$  are two real numbers between 0 and 1 and  $M = \langle S, A, T, \mathcal{R}, \gamma \rangle$  is any MDP. There exists an input  $m = m(\frac{1}{\epsilon}, \frac{1}{\delta})$ , satisfying  $m(\frac{1}{\epsilon}, \frac{1}{\delta}) = O(\frac{|S|}{\epsilon^2(1-\gamma)^4} + \frac{1}{\epsilon^2(1-\gamma)^4} \ln \frac{|S||A|}{\epsilon(1-\gamma)^\delta})$ , and  $\beta = (1/(1-\gamma))\sqrt{\ln(2|S||A|m/\delta)}/2$  such that if MBIE-EB is executed on MDP  $M$ , then the following holds. Let  $\mathcal{A}_t$  denote MBIE-EB's policy at time  $t$  and  $s_t$  denote the state at time  $t$ . With probability at least  $1 - \delta$ ,  $V_M^{\mathcal{A}_t}(s_t) \geq V_M^*(s_t) - \epsilon$  is true for all but  $O(\frac{|S||A|}{\epsilon^3(1-\gamma)^6}(|S| + \ln \frac{|S||A|}{\epsilon(1-\gamma)^\delta}) \ln \frac{1}{\delta} \ln \frac{1}{\epsilon(1-\gamma)})$  timesteps  $t$ .

# A Sufficient Set of Conditions to Make a RL Algorithm PAC

- Strehl, A. L., Li, L., & Littman, M. L. (2006). Incremental model-based learners with formal learning-time guarantees. In Proceedings of the Twenty-Second Conference on Uncertainty in Artificial Intelligence (pp. 485-493)

# A Sufficient Set of Conditions to Make a RL Algorithm PAC





# How Does MBIE-EB Fulfill these Conditions?





# Table of Contents

1 MDPs

2 Bayesian MDPs

3 Generalization and Exploration

4 Summary

# Refresher: Bayesian Bandits

- **Bayesian bandits** exploit prior knowledge of rewards,  $p[\mathcal{R}]$
- They compute posterior distribution of rewards  $p[\mathcal{R} \mid h_t]$ , where  $h_t = (a_1, r_1, \dots, a_{t-1}, r_{t-1})$
- Use posterior to guide exploration
  - Upper confidence bounds (Bayesian UCB)
  - Probability matching (Thompson Sampling)
- Better performance if prior knowledge is accurate

# Refresher: Bernoulli Bandits

- Consider a bandit problem where the reward of an arm is a binary outcome  $\{0, 1\}$  sampled from a Bernoulli with parameter  $\theta$ 
  - E.g. Advertisement click through rate, patient treatment succeeds/fails, ...
- The Beta distribution  $Beta(\alpha, \beta)$  is conjugate for the Bernoulli distribution

$$p(\theta|\alpha, \beta) = \theta^{\alpha-1}(1-\theta)^{\beta-1} \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)}$$

where  $\Gamma(x)$  is the Gamma function.

- Assume the prior over  $\theta$  is a  $Beta(\alpha, \beta)$  as above
- Then after observed a reward  $r \in \{0, 1\}$  then updated posterior over  $\theta$  is  $Beta(r + \alpha, 1 - r + \beta)$

# Thompson Sampling for Bandits

- 
- 1: Initialize prior over each arm  $a$ ,  $p(\mathcal{R}_a)$
  - 2: **loop**
  - 3:   For each arm  $a$  **sample** a reward distribution  $\mathcal{R}_a$  from posterior
  - 4:   Compute action-value function  $Q(a) = \mathbb{E}[\mathcal{R}_a]$
  - 5:    $a_t = \arg \max_{a \in \mathcal{A}} Q(a)$
  - 6:   Observe reward  $r$
  - 7:   Update posterior  $p(\mathcal{R}_a|r)$  using Bayes law
  - 8: **end loop**
-

# Bayesian Model-Based RL

- Maintain posterior distribution over **MDP** models
- Estimate both transition and rewards,  $p[\mathcal{P}, \mathcal{R} \mid h_t]$ , where  $h_t = (s_1, a_1, r_1, \dots, s_t)$  is the history
- Use posterior to guide exploration
  - Upper confidence bounds (Bayesian UCB)
  - Probability matching (Thompson sampling)

# Thompson Sampling: Model-Based RL

- Thompson sampling implements probability matching

$$\begin{aligned}\pi(s, a \mid h_t) &= \mathbb{P}[Q(s, a) \geq Q(s, a'), \forall a' \neq a \mid h_t] \\ &= \mathbb{E}_{\mathcal{P}, \mathcal{R} \mid h_t} \left[ \mathbb{1}(a = \arg \max_{a \in \mathcal{A}} Q(s, a)) \right]\end{aligned}$$

- Use Bayes law to compute posterior distribution  $p[\mathcal{P}, \mathcal{R} \mid h_t]$
- Sample** an MDP  $\mathcal{P}, \mathcal{R}$  from posterior
- Solve MDP using favorite planning algorithm to get  $Q^*(s, a)$
- Select optimal action for sample MDP,  $a_t = \arg \max_{a \in \mathcal{A}} Q^*(s_t, a)$

# Thompson Sampling for MDPs

- 
- 1: Initialize prior over the dynamics and reward models for each  $(s, a)$ ,  $p(\mathcal{R}_{as}), p(\mathcal{T}(s'|s, a))$
  - 2: Initialize state  $s_0$
  - 3: **loop**
  - 4:   Sample a MDP  $\mathcal{M}$ : for each  $(s, a)$  pair, sample a dynamics model  $\mathcal{T}(s'|s, a)$  and reward model  $\mathcal{R}(s, a)$
  - 5:   Compute  $Q_{\mathcal{M}}^*$ , optimal value for MDP  $\mathcal{M}$
  - 6:    $a_t = \arg \max_{a \in \mathcal{A}} Q_{\mathcal{M}}^*(s_t, a)$
  - 7:   Observe reward  $r_t$  and next state  $s_{t+1}$
  - 8:   Update posterior  $p(\mathcal{R}_{a_t s_t} | r_t), p(\mathcal{T}(s' | s_t, a_t) | s_{t+1})$  using Bayes rule
  - 9:    $t = t + 1$
  - 10: **end loop**
-

# Check Your Understanding: Fast RL III

- Strategic exploration in MDPs (select all):
  - 1 Doesn't really matter because the distribution of data is independent of the policy followed
  - 2 Can involve using optimism with respect to both the possible dynamics and reward models in order to compute an optimistic Q function
  - 3 Is known as PAC if the number of time steps on which a less than near optimal decision is made is guaranteed to be less than an exponential function of the problem domain parameters (state space cardinality, etc).
  - 4 Not sure
- In Thompson sampling for MDPs:
  - 1 TS samples the reward model parameters and could use the empirical average for the dynamics model parameters and obtain the same performance
  - 2 Must perform MDP planning everytime the posterior is updated
  - 3 Has the same computational cost each step as Q-learning
  - 4 Not sure



# Resampling in Coordinated Exploration

- Concurrent PAC RL. Guo and Brunskill. AAAI 2015
- Coordinated Exploration in Concurrent Reinforcement Learning. Dimakopoulou and Van Roy. ICML 2018
- <https://www.youtube.com/watch?v=xjGK-wm0Pkl&feature=youtu.be>

# Table of Contents

- 1 MDPs
- 2 Bayesian MDPs
- 3 Generalization and Exploration**
- 4 Summary

# Generalization and Strategic Exploration

- Active area of ongoing research: combine generalization & strategic exploration
- Many approaches are grounded by principles outlined here
  - Optimism under uncertainty
  - Thompson sampling

# Generalization and Optimism

- Recall MBIE-EB algorithm for finite state and action domains
- What needs to be modified for continuous / extremely large state and/or action spaces?

# Model-Based Interval Estimation with Exploration Bonus (MBIE-EB)

(Strehl and Littman, J of Computer & Sciences 2008)

- 
- 1: Given  $\epsilon, \delta, m$
  - 2:  $\beta = \frac{1}{1-\gamma} \sqrt{0.5 \ln(2|S||A|m/\delta)}$
  - 3:  $n_{sas}(s, a, s') = 0, \forall s \in S, a \in A, s' \in S$
  - 4:  $rc(s, a) = 0, n_{sa}(s, a) = 0, \tilde{Q}(s, a) = 1/(1 - \gamma), \forall s \in S, a \in A$
  - 5:  $t = 0, s_t = s_{init}$
  - 6: **loop**
  - 7:    $a_t = \arg \max_{a \in A} \tilde{Q}(s_t, a)$
  - 8:   Observe reward  $r_t$  and state  $s_{t+1}$
  - 9:    $n_{sa}(s_t, a_t) = n_{sa}(s_t, a_t) + 1, n_{sas}(s_t, a_t, s_{t+1}) = n_{sas}(s_t, a_t, s_{t+1}) + 1$
  - 10:    $rc(s_t, a_t) = \frac{rc(s_t, a_t)(n_{sa}(s_t, a_t) - 1) + r_t}{n_{sa}(s_t, a_t)}$
  - 11:    $\hat{R}(s_t, a_t) = rc(s_t, a_t)$  and  $\hat{T}(s'|s_t, a_t) = \frac{n_{sas}(s_t, a_t, s')}{n_{sa}(s_t, a_t)}, \forall s' \in S$
  - 12:   **while** not converged **do**
  - 13:      $\tilde{Q}(s, a) = \hat{R}(s, a) + \gamma \sum_{s'} \hat{T}(s'|s, a) \max_{a'} \tilde{Q}(s', a) + \frac{\beta}{\sqrt{n_{sa}(s, a)}}, \forall s \in S, a \in A$
  - 14:   **end while**
  - 15: **end loop**

# Generalization and Optimism

- Recall MBIE-EB algorithm for finite state and action domains
- What needs to be modified for continuous / extremely large state and/or action spaces?
- Estimating uncertainty
  - Counts of  $(s,a)$  and  $(s,a,s')$  tuples are not useful if we expect only to encounter any state once
- Computing a policy
  - Model-based planning will fail
- So far, model-free approaches have generally had more success than model-based approaches for extremely large domains
  - Building good transition models to predict pixels is challenging

# Recall: Value Function Approximation with Control

- For Q-learning use a TD target  $r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w})$  which leverages the max of the current function approximation value

$$\Delta \mathbf{w} = \alpha(r(s) + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

# Recall: Value Function Approximation with Control

- For Q-learning use a TD target  $r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w})$  which leverages the max of the current function approximation value

$$\Delta \mathbf{w} = \alpha (r(s) + r_{bonus}(s, a) + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$



# Recall: Value Function Approximation with Control

- For Q-learning use a TD target  $r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w})$  which leverages the max of the current function approximation value

$$\Delta \mathbf{w} = \alpha (r(s) + r_{bonus}(s, a) + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

- $r_{bonus}(s, a)$  should reflect uncertainty about future reward from  $(s, a)$
- Approaches for deep RL that make an estimate of visits / density of visits include: Bellemare et al. NIPS 2016; Ostrovski et al. ICML 2017; Tang et al. NIPS 2017
- Note: bonus terms are computed at time of visit. During episodic replay can become outdated.

# Benefits of Strategic Exploration: Montezuma's revenge




Figure 3: “Known world” of a DQN agent trained for 50 million frames with (**right**) and without (**left**) count-based exploration bonuses, in MONTEZUMA’S REVENGE.

**Figure:** Bellemare et al. "Unifying Count-Based Exploration and Intrinsic Motivation"

- Enormously better than standard DQN with  $\epsilon$ -greedy approach

# Generalization and Strategic Exploration: Thompson Sampling

- Leveraging Bayesian perspective has also inspired some approaches
- One approach: Thompson sampling over representation & parameters (Mandel, Liu, Brunskill, Popovic IJCAI 2016)

# Generalization and Strategic Exploration: Thompson Sampling

- Leveraging Bayesian perspective has also inspired some approaches
- One approach: Thompson sampling over representation & parameters (Mandel, Liu, Brunskill, Popovic IJCAI 2016)
- For scaling up to very large domains, again useful to consider model-free approaches
- Non-trivial: would like to be able to sample from a posterior over possible  $Q^*$
- Bootstrapped DQN (Osband et al. NIPS 2016)
  - Train  $C$  DQN agents using bootstrapped samples
  - When acting, choose action with highest  $Q$  value over any of the  $C$  agents
  - Some performance gain, not as effective as reward bonus approaches

# Generalization and Strategic Exploration: Thompson Sampling

- Leveraging Bayesian perspective has also inspired some approaches
- One approach: Thompson sampling over representation & parameters (Mandel, Liu, Brunskill, Popovic IJCAI 2016)
- For scaling up to very large domains, again useful to consider model-free approaches
- Non-trivial: would like to be able to sample from a posterior over possible  $Q^*$
- Bootstrapped DQN (Osband et al. NIPS 2016)
- Efficient Exploration through Bayesian Deep Q-Networks (Azzaddenesheli, Anandkumar, NeurIPS workshop 2017)
  - Use deep neural network
  - On last layer use Bayesian linear regression
  - Be optimistic with respect to the resulting posterior
  - Very simple, empirically much better than just doing linear regression on last layer or bootstrapped DQN, not as good as reward bonuses in some cases

# Table of Contents

- 1 MDPs
- 2 Bayesian MDPs
- 3 Generalization and Exploration
- 4 Summary

# Summary: What You Are Expected to Know

- Define the tension of exploration and exploitation in RL and why this does not arise in supervised or unsupervised learning
- Be able to define and compare different criteria for "good" performance (empirical, convergence, asymptotic, regret, PAC)
- Be able to map algorithms discussed in detail in class to the performance criteria they satisfy

# Class Structure

- Last time: Fast Learning (Bayesian bandits to MDPs)
- **This time: Fast Learning III (MDPs)**
- Next time: Batch RL