

RL: Deep

Double DQN

Marius Lindauer



Winter Term 2021

Recall: Double Q-Learning

► Initialization:

- $Q_1(s, a)$ and $Q_2(s, a)$ for $\forall s \in S, a \in A$
- $t = 0$
- initial state $s_t = s_0$

► Loop

- Select a_t using ϵ -greedy $\pi(s) \in \arg \max_{a \in A} Q_1(s_t, a) + Q_2(s_t, a)$
- Observe (r_t, s_{t+1})
- With 50-50 probability either
 1. $Q_1(s_t, a_t) \leftarrow Q_1(s_t, a_t) + \alpha(r_t + \gamma \max_{a \in A} Q_2(s_{t+1}, a) - Q_1(s_t, a_t))$
or
 2. $Q_2(s_t, a_t) \leftarrow Q_2(s_t, a_t) + \alpha(r_t + \gamma \max_{a \in A} Q_1(s_{t+1}, a) - Q_2(s_t, a_t))$
- $t = t + 1$

Recall: Double Q-Learning

► Initialization:

- $Q_1(s, a)$ and $Q_2(s, a)$ for $\forall s \in S, a \in A$
- $t = 0$
- initial state $s_t = s_0$

► Loop

- Select a_t using ϵ -greedy $\pi(s) \in \arg \max_{a \in A} Q_1(s_t, a) + Q_2(s_t, a)$
- Observe (r_t, s_{t+1})
- With 50-50 probability either
 1. $Q_1(s_t, a_t) \leftarrow Q_1(s_t, a_t) + \alpha(r_t + \gamma \max_{a \in A} Q_2(s_{t+1}, a) - Q_1(s_t, a_t))$
or
 2. $Q_2(s_t, a_t) \leftarrow Q_2(s_t, a_t) + \alpha(r_t + \gamma \max_{a \in A} Q_1(s_{t+1}, a) - Q_2(s_t, a_t))$
- $t = t + 1$

~> reduces maximization bias

Double DQN Hasselt et al. 2015

- ▶ Extend this idea to DQN
- ▶ Current Q-network \vec{w} is used to select actions
- ▶ Older Q-network \vec{w}^- is used to evaluate actions
- ▶ TD-error:

$$r + \gamma \underbrace{\hat{Q}(s', \arg \max_{a' \in A} \hat{Q}(s', a'; \vec{w}); \vec{w}^-)}_{\text{Action selection: } \vec{w}} - Q(s, a; \vec{w})$$

Action evaluation: \vec{w}^-

Double DQN Hasselt et al. 2015

- ▶ Extend this idea to DQN
- ▶ Current Q-network \vec{w} is used to select actions
- ▶ Older Q-network \vec{w}^- is used to evaluate actions
- ▶ TD-error:

$$r + \gamma \underbrace{\hat{Q}(s', \arg \max_{a' \in A} \hat{Q}(s', a'; \vec{w}); \vec{w}^-)}_{\text{Action evaluation: } \vec{w}^-} - \underbrace{Q(s, a; \vec{w})}_{\text{Action selection: } \vec{w}}$$

- ▶ Allows flipping between both weight sets frequently
 - ▶ alternatively, Polyak averaging:

$$w' \leftarrow \tau w + (1 - \tau)w'$$

- ▶ τ is fairly small, e.g. 0.01
- ▶ Faster propagation of information compared to original DQN

Clipped Double DQN Fujimoto et al. 2018

- ▶ Extend this idea to DQN
- ▶ Again having two independent Q-networks with \vec{w}_1 and \vec{w}_2
- ▶ Take minimum action value for successor state
- ▶ TD-error:

$$r + \gamma \min_{i=\{1,2\}} Q(s', \arg \max_{a' \in A} Q(s', a'; \vec{w}_i); \vec{w}_i) - Q(s, a; \vec{w})$$

- ▶ Less overestimation of Q-values
- ▶ More stable learning targets