

# RL: Policy Search

## Gradient-free Optimization

Marius Lindauer



# Policy optimization

- Policy based reinforcement learning is an **optimization** problem over  $\theta$
- ↪ Find policy parameters  $\theta^*$  that maximize  $V(s_0, \theta^*)$
- We can use gradient-free approaches (a.k.a. black-box optimization)
  - ▶ Hill climbing
  - ▶ Simplex / amoeba / Nelder Mead
  - ▶ Genetic algorithms
  - ▶ Cross-Entropy method
  - ▶ Covariance Matrix Adaptation (CMA)

# Policy optimization


- Policy based reinforcement learning is an **optimization** problem over  $\theta$ 
  - ↪ Find policy parameters  $\theta^*$  that maximize  $V(s_0, \theta^*)$
- We can use gradient-free approaches (a.k.a. black-box optimization)
  - ▶ Hill climbing
  - ▶ Simplex / amoeba / Nelder Mead
  - ▶ Genetic algorithms
  - ▶ Cross-Entropy method
  - ▶ Covariance Matrix Adaptation (CMA)
- gradient-free optimizers are (often) designed for
  - ▶ many function evaluations → possible in RL
  - ▶ parallel computation → possible in RL
  - ▶ a few to hundreds of dimensions → RL?

# Policy optimization

- Policy based reinforcement learning is an **optimization** problem over  $\theta$
- ↪ Find policy parameters  $\theta^*$  that maximize  $V(s_0, \theta^*)$
- We can use gradient-free approaches (a.k.a. black-box optimization)
  - ▶ Hill climbing
  - ▶ Simplex / amoeba / Nelder Mead
  - ▶ Genetic algorithms
  - ▶ Cross-Entropy method
  - ▶ Covariance Matrix Adaptation (CMA)
- gradient-free optimizers are (often) designed for
  - ▶ many function evaluations → possible in RL
  - ▶ parallel computation → possible in RL
  - ▶ a few to hundreds of dimensions → RL?
- if we encode the policy  $\pi_\theta$  as a DNN, we might have **millions** of dimensions (i.e., parameters in  $\theta$ )

# Policy Optimization with Evolutionary Strategies

- Evolutionary Strategies can perform surprisingly well nevertheless  
[Salimans et al. 2017; Chrabaszcz et al. 2018; Fuks et al. 2019]



Sources: [Chrabaszcz et al. 2018; Fuks et al. 2019]

# Policy Optimization with Evolutionary Strategies

[Chrabaszcz et al. 2018]

---

## Algorithm 1: Canonical Evolution Strategy

---

**Input:**

$\theta_0$  - Initial policy vector parameters

$T$  - time budget

$\lambda$  - Population size

$\mu$  - Parent population size

$\sigma$  - Mutation step-size

$F(\theta)$  - Fitness function for policy evaluation

```
1 for  $j \in \{1 \dots \mu\}$  do
2    $w_j = \frac{\log(\mu+0.5) - \log(j)}{\sum_{k=1}^{\mu} \log(\mu+0.5) - \log(k)}$ 
3 for  $t = 0, 1, \dots, T$  do
4   for  $i = 1, 2, \dots, \lambda$  do
5      $\epsilon_i \sim \mathcal{N}(0, I)$ 
6      $s_i \leftarrow F(\theta_t + \sigma \cdot \epsilon_i)$ 
7   Sort  $(\epsilon_1, \dots, \epsilon_\lambda)$  according to  $s$  in ascending order
8   Update policy:  $\theta_{t+1} \leftarrow \theta_t + \sigma \cdot \sum_{j=1}^{\mu} w_j \cdot \epsilon_j$ 
```

**Output:** Return best found policy  $\theta_t$

---

# Progressive Episode Length [Fuks et al. 2019]

