

# RL: Policy Search

## Finite Difference

Marius Lindauer



# Policy Gradient

- Assume episodic MDPs
- Policy gradient algorithms search for a **local** maximum in  $V(s_0, \theta)$  by ascending the gradient of the policy, w.r.t parameters  $\theta$

$$\Delta\theta = \alpha \nabla_{\theta} V(s_0, \theta)$$

where  $\alpha$  is the learning rate (step-size) and  $\nabla_{\theta} V(s_0, \theta)$  is the policy gradient

$$\nabla_{\theta} V(s_0, \theta) = \begin{pmatrix} \frac{\partial V(s_0, \theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial V(s_0, \theta)}{\partial \theta_n} \end{pmatrix}$$

# Simple Approach: Compute Gradients by Finite Differences

- To evaluate policy gradient of  $\pi_{\theta}(s, a)$
- For each dimension  $k \in [1, n]$ 
  - ▶ Estimate  $k$ -th partial derivative of objective function wrt  $\theta$
  - ▶ By perturbing  $\theta$  by small amount  $\epsilon$  in  $k$ -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon u_k) - V(s_0, \theta)}{\epsilon}$$

where  $u_k$  is a unit vector with 1 in  $k$ -th component and 0 elsewhere

# Simple Approach: Compute Gradients by Finite Differences

- To evaluate policy gradient of  $\pi_{\theta}(s, a)$
- For each dimension  $k \in [1, n]$ 
  - ▶ Estimate  $k$ -th partial derivative of objective function wrt  $\theta$
  - ▶ By perturbing  $\theta$  by small amount  $\epsilon$  in  $k$ -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon u_k) - V(s_0, \theta)}{\epsilon}$$

where  $u_k$  is a unit vector with 1 in  $k$ -th component and 0 elsewhere

- $\epsilon$  should be
  - ▶ large enough to observe a change in the policy and value function
  - ▶ small enough to have a good gradient approximation

# Simple Approach: Compute Gradients by Finite Differences

- To evaluate policy gradient of  $\pi_{\theta}(s, a)$
- For each dimension  $k \in [1, n]$ 
  - ▶ Estimate  $k$ -th partial derivative of objective function wrt  $\theta$
  - ▶ By perturbing  $\theta$  by small amount  $\epsilon$  in  $k$ -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon u_k) - V(s_0, \theta)}{\epsilon}$$

where  $u_k$  is a unit vector with 1 in  $k$ -th component and 0 elsewhere

- $\epsilon$  should be
  - ▶ large enough to observe a change in the policy and value function
  - ▶ small enough to have a good gradient approximation
- Uses  $\geq n$  evaluations to compute policy gradient in  $n$  dimensions
  - ~ fairly inefficient for doing a single update!
  - ~ weight space should be small — no computer vision

# Simple Approach: Compute Gradients by Finite Differences

- To evaluate policy gradient of  $\pi_{\theta}(s, a)$
- For each dimension  $k \in [1, n]$ 
  - ▶ Estimate  $k$ -th partial derivative of objective function wrt  $\theta$
  - ▶ By perturbing  $\theta$  by small amount  $\epsilon$  in  $k$ -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon u_k) - V(s_0, \theta)}{\epsilon}$$

where  $u_k$  is a unit vector with 1 in  $k$ -th component and 0 elsewhere

- $\epsilon$  should be
  - ▶ large enough to observe a change in the policy and value function
  - ▶ small enough to have a good gradient approximation
- Uses  $\geq n$  evaluations to compute policy gradient in  $n$  dimensions
  - ~ fairly inefficient for doing a single update!
  - ~ weight space should be small — no computer vision
- Simple, noisy, inefficient – but sometimes effective
- Works for arbitrary policies, even if policy is not differentiable