

Meta Reinforcement Learning

The Big Picture

Marius Lindauer



Can We Generalize beyond a given MDP?

- What happens if the environment changes? (non-stationary environments)
 - ▶ Can we efficiently adapt our policy to changed transitions or reward functions?
 - After a human player learned how to play Super Mario Bros in the first levels, they will also play fairly well the upcoming levels.
 - However, an RL agent potentially will fail.
- ~> Strong limitations regarding the applications of a trained agent

- 1 Can we use easy environments to learn how to behave in hard environments?

Main Ideas

- 1 Can we use easy environments to learn how to behave in hard environments?
- 2 Can we train a policy that is easily adaptable to new environments?

Main Ideas

- 1 Can we use easy environments to learn how to behave in hard environments?
- 2 Can we train a policy that is easily adaptable to new environments?
- 3 Can we find better training dynamics across a set of environments?

Main Ideas

- ① Can we use easy environments to learn how to behave in hard environments?
- ② Can we train a policy that is easily adaptable to new environments?
- ③ Can we find better training dynamics across a set of environments?
- ④ Can we train a policy that generalizes to new environments without any new training?
- Assumption: We sample our environments i.i.d. from a fixed distribution
 - ▶ Similar to the assumption in supervised learning, but on a meta-level
 - ▶ Training environments to train our agent on and test environments to check how well it performs.
 - ▶ We might have control how we sample from this distribution; we might don't.