

Curriculum Reinforcement Learning

SPaCE

Marius Lindauer



Winter Term 2021

- ▶ Most agents estimate the state-value function V in some way
- ▶ Common definition of V :

$$V(s) = \mathbb{E}_{\pi} \sum_t^T \gamma^t \cdot r_t$$

- ▶ Therefore $V(s_0)$ estimates the total discounted reward for the whole episode
- ▶ Idea: use this information for curriculum generation

Setting – what makes SPaCE special?

- ▶ Designed for deep contextual RL
- ▶ No prior knowledge about instance space required, e.g.:
 - ▶ size
 - ▶ difficulty
 - ▶ difficulty regions

Algorithm Outline

- ▶ Hyperparameters: threshold η , increment size κ
- ▶ Until desired number of steps is reached:
 - ▶ Choose the instances on which the evaluation has changed most (according to current instance set size)
 - ▶ Train on those instances
 - ▶ Evaluate if performance on the training set has changed by at least η . If not, increase instance set size by κ

Results - AntGoal




Figure: Comparing SPaCE and round robin on AntGoal with broken limbs.

Results - PointMass




Figure: Comparing SPaCE and round robin on contextual PointMass.